

Daniel Iván Parra Verde

AI/ML Engineer • AWS Certified 2x (AIP-C01, SAA-C03) • AI Agents & ML Platforms • +4 years
 Portfolio: danivpv.com | Blog: danivpv.com/blog | GitHub: [danivpv](https://github.com/danivpv) | LinkedIn: [danivpv](https://www.linkedin.com/in/danivpv)
 Tel. +52 844 783 6082 | danivpv@outlook.com | Saltillo, México

EDUCATION

UNAM-CFATA

2020 | BSc. TECHNOLOGY

GRONINGEN UNIVERSITY (G)

2019 | EXCHANGE (AI/CS)

Netherlands; Full Scholarship

UNAM-IIMAS

2022 | MSc. CANDIDATE IN MATH

SKILLS

AI and ML Libraries:

PyTorch • LangChain • LlamaIndex •
 Hugging Face • scikit-learn • Pandas •
 NumPy • Docling • Gemini • OpenAI •
 Anthropic • OpenCV

Frameworks:

FastAPI • Temporal • Spark • Next.js •
 SQLAlchemy • Django • Kafka

Databases:

PostgreSQL • MongoDB • Redis •
 Qdrant • SQLite

Programming Languages:

Python • SQL • Bash • TypeScript •
 Terraform • Julia • Wolfram

Software Tools:

Docker • Git • GitHub Actions • Sentry
 • Postman • $\LaTeX$

Spoken Languages

English (+C1) • German (+A2) • Spanish

CERTIFICATIONS

AWS GenAI Developer Professional

AWS Solutions Architect Associate

COURSEWORK

Bachelor

Deep Learning in **Python**(G)
 Relational Databases in **SQL** (G)
 Artificial Intelligence in **Python** (G)
 Algorithms and Data Structures in **C** (G)
 Linear Algebra

Master

Statistical Inference
 Theoretic Machine Learning
 Information Theory
 Statistical Foundations of Privacy
 Numerical Analysis **Python**, **Julia**
 Real and Functional Analysis

WORK EXPERIENCE

SENIOR DATA SCIENTIST

Arkham Technologies | April 2025 - June 2026 | Remote

- Deployed AI models into production and automated AI agent monitoring (LLM-as-Judge) via **FastAPI** microservices, **PostgreSQL** and **AWS**.
- Shipped cross-provider LLM integrations serving hundreds of users per week.
- Engineered enterprise Machine Learning pipelines; scaled 1K models on **AWS EMR** via **Spark** and event-driven architecture (**Kafka**, **Temporal**).
- Led **MLOps** & **CI/CD** of **Docker** images for prod ML platform via **GitHub Actions**, **AWS** and **Terraform**. 10× faster builds vs baseline.
- Debugged **FastAPI/Kafka** ML and AI microservices via **CloudWatch** and **Sentry**.

AI MENTOR & DEVELOPER

LIDR | March 2025 - Present | Remote

- Deliver live-coding on AI-assisted development, **Cursor**, **MCP**, **OpenRouter**, integrations, and developer tooling.

DATA SCIENTIST

Kuona | November 2023 - June 2024 | Remote

- Designed QLoRA and sales-weighted loss functions for critical revenue KPIs.
- Tuned **PyTorch RNNs** for enterprise sales forecasting, turning predictions into actionable decision intelligence for retail revenue management (12% MAPE).
- Modeled competitors' pricing data, reducing forecasting MAPE by 40% and enabling price-demand elasticity analysis for business decisions via **Pandas**.

DATA SCIENTIST JR.

Entropia AI | February 2023 - November 2023 | Remote

- Deployed **RAG** AI solutions into production; raised custom QA LLM-as-Judge accuracy from 8% to 72% on legal corpora using **LlamaIndex** and **MongoDB**.
- Built 8 **ETL** pipelines (**Python**, **bs4**, **Airflow**) ingesting 2K docs with 100M+ tokens from official Mexican government sites into **MongoDB** for legal partners.

SOFTWARE DEVELOPER

Wolfram Research | February 2021 - February 2022 | Remote

- Sysadmin. Fullstack dev of portal for "Wolfram Summer Camp" (+40 users).
- Published paper "Turing Patterns in Networks" in Complex Systems Journal.

PROJECTS

AUTONOMOUS FINOPS AGENT

Live Demo | GitHub Repository

- Architected AWS native agent with tools via **Bedrock AgentCore Gateway** and versioned prompts with **Bedrock Prompt Management**.
- Built full-stack AI SaaS deployed with **AWS CDK**, **Next.js** on **Amplify**, and serverless backend for cloud economics analytics.

ENTERPRISE AUTOML PLATFORM

Architecture Deep Dive | GitHub Repository

- Architected 4 core subsystems via **AWS CDK**: feature store (**Feast**), lineage (**MLflow**), catalog API, and decoupled compute for inference and training.
- Enforced VPC network isolation with private subnets, **ALB** routing, and **NAT** egress for secure **MLflow** and batch inference.